

Practical Inference

EN.705.743: ChatGPT from Scratch

Outline

- Recap prompting ideas
- LLM inference as a subroutine
- RAG
- Popular Inference Providers + Libraries
- Huggingface Overview
 - Inference API
 - Datasets
 - Transformers

Prompting Recap

Last time we saw that introducing a pattern to the model can induce a specific behavior:

I love the beach / Ich liebe den Strand

What time is it? / Wie spät ist es?

The beer is ten dollars. / Das Bier kostet zehn Dollar.

Translation examples.
Clear English / German
pattern.

Where can I get a coffee? /

← Incomplete pattern continuation
for what we want to translate.

Prompting Recap


Last time we saw that introducing a pattern to the model can induce a specific behavior:

I love the beach / Ich liebe den Strand

What time is it? / Wie spät ist es?

The beer is ten dollars. / Das Bier kostet zehn Dollar.

Where can I get a coffee? / Wo kann ich einen Kaffee bekommen?



By continuing the text, the model serves as a translator.

Prompting + LLM as Subroutine

Taking this further, we can store a variety of these prompts to call forth a given behavior as needed. These can replace or work alongside traditional software modules.

Prompting + LLM as Subroutine

Taking this further, we can store a variety of these prompts to call forth a given behavior as needed. These can replace or work alongside traditional software modules.

```
def convert_to_F(cel):  
    return round(cel*(9.0/5.0) + 32.0, 1)  
  
def city_uses_fahrenheit(city):  
    # not so easy...  
    pass  
  
def format_weather_reply(city, temp):  
    if city_uses_fahrenheit(city):  
        return f"The temperature is {convert_to_F(temp)} F."  
    else:  
        return f"The temperature is {temp} C."
```

Prompting + LLM as Subroutine

Taking this further, we can store a variety of these prompts to call forth a given behavior as needed. These can replace or work alongside traditional software modules.

In the example, we have replaced a typical method with a call to an LLM, and we are trusting that the LLM knows which cities use which system.

Note that pre/post processing and error handling is important.

```
def convert_to_F(cel):
    return round(cel*(9.0/5.0) + 32.0, 1)

def city_uses_fahrenheit(city):
    prompt = """
    Chicago: Chicago is in the US which uses Farenheit. True.
    Paris: Paris is in France which uses Celsius. False.
    New York: New York is in the US which uses Farenheit. True.
    London: London is in the UK which uses Celsius. False.
    Hong Kong: Hong Kong is in China which uses Celsius. False.
    """

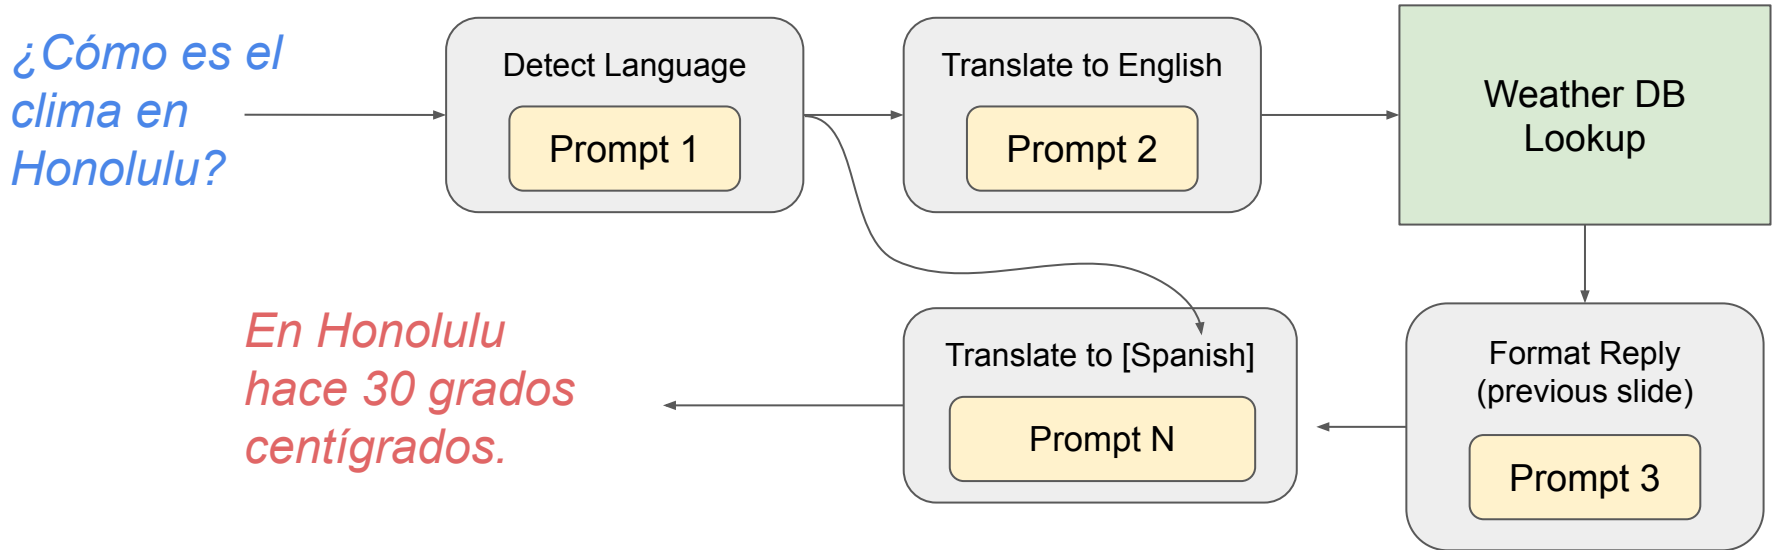
    query = prompt+"\n"+city+": "
    output = LLM(query)

    return "true" in output.lower()

def format_weather_reply(city, temp):
    if city_uses_fahrenheit(city):
        return f"The temperature is {convert_to_F(temp)} F."
    else:
        return f"The temperature is {temp} C."
```

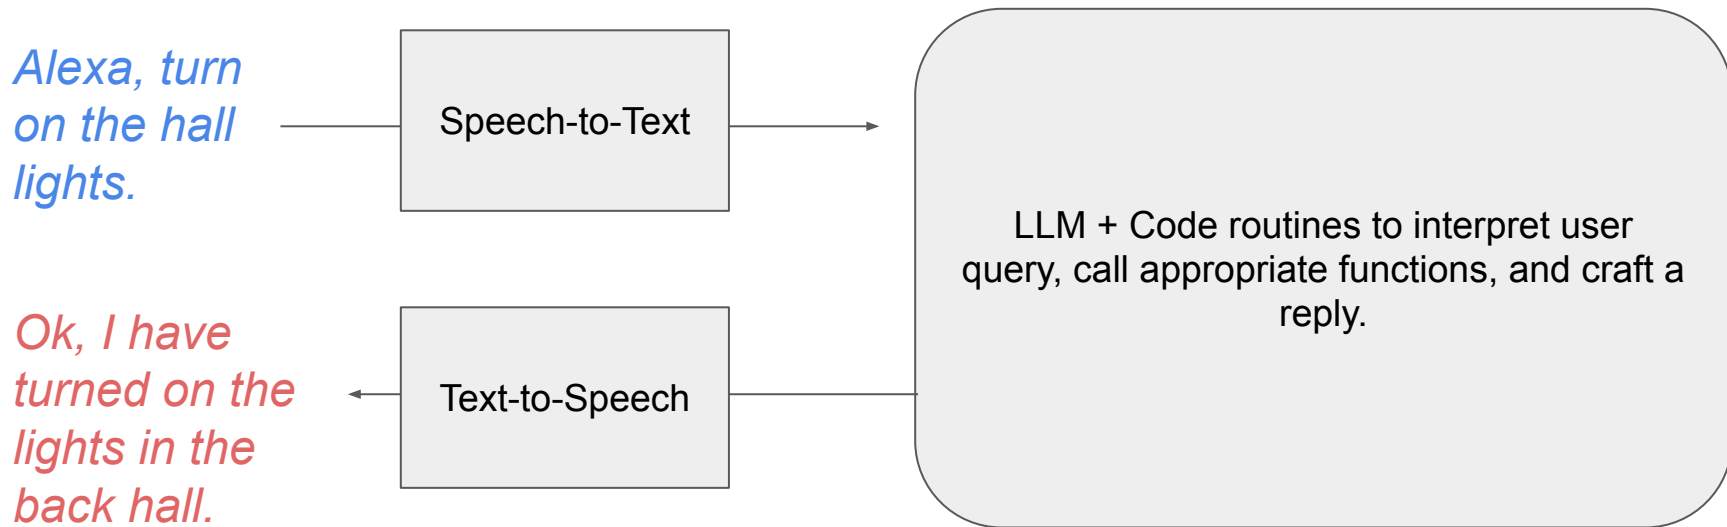
Prompting + LLM as Subroutine

We could string together a variety of these methods to create complex software systems. Since our own code is still orchestrating everything, we can route prompts and responses as needed.



Conversational Interfaces

This is especially useful when creating a system that will use a conversational interface, for example a voice-controlled app.



Retrieval Augmented Generation (RAG)

Retrieval Augmented Generation (RAG)

What if we need to include specialized data in the LLM query? In the prior examples, the LLM itself would not know the current weather or the state of the hallway lighting.

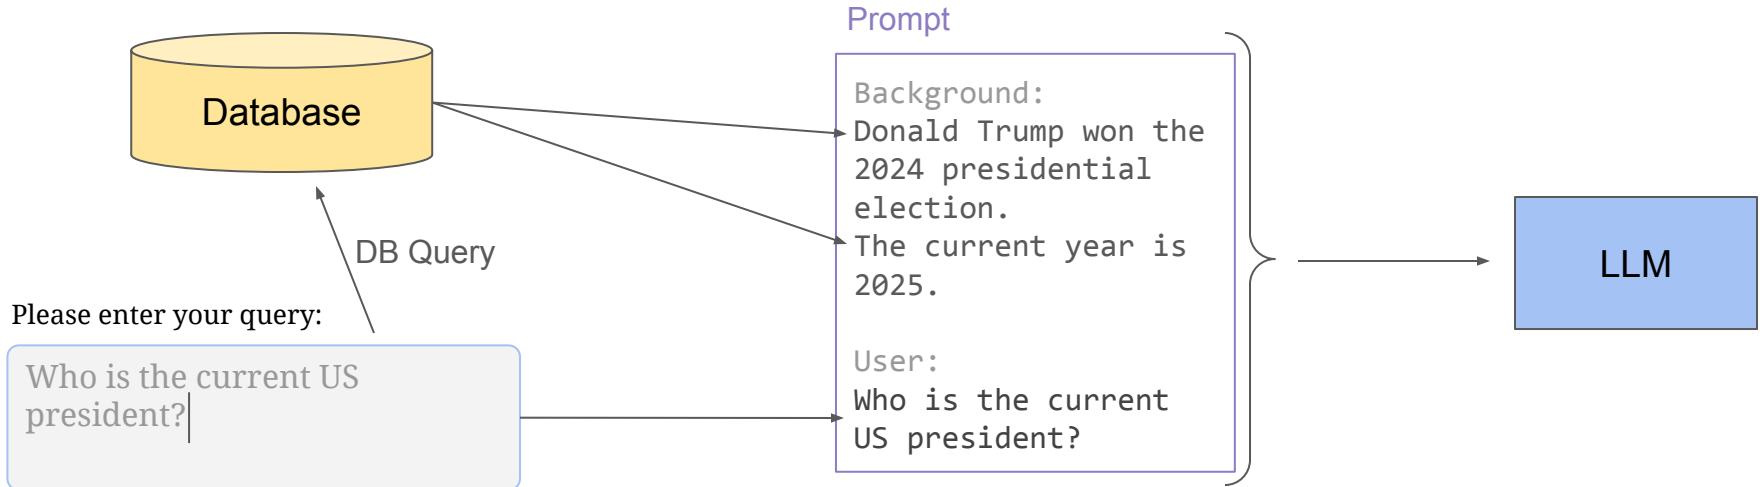
More broadly, an LLM will not have access to live information or non-public data sources.

To solve this, we can inject information from another data source into our prompts.

Retrieval Augmented Generation (RAG)

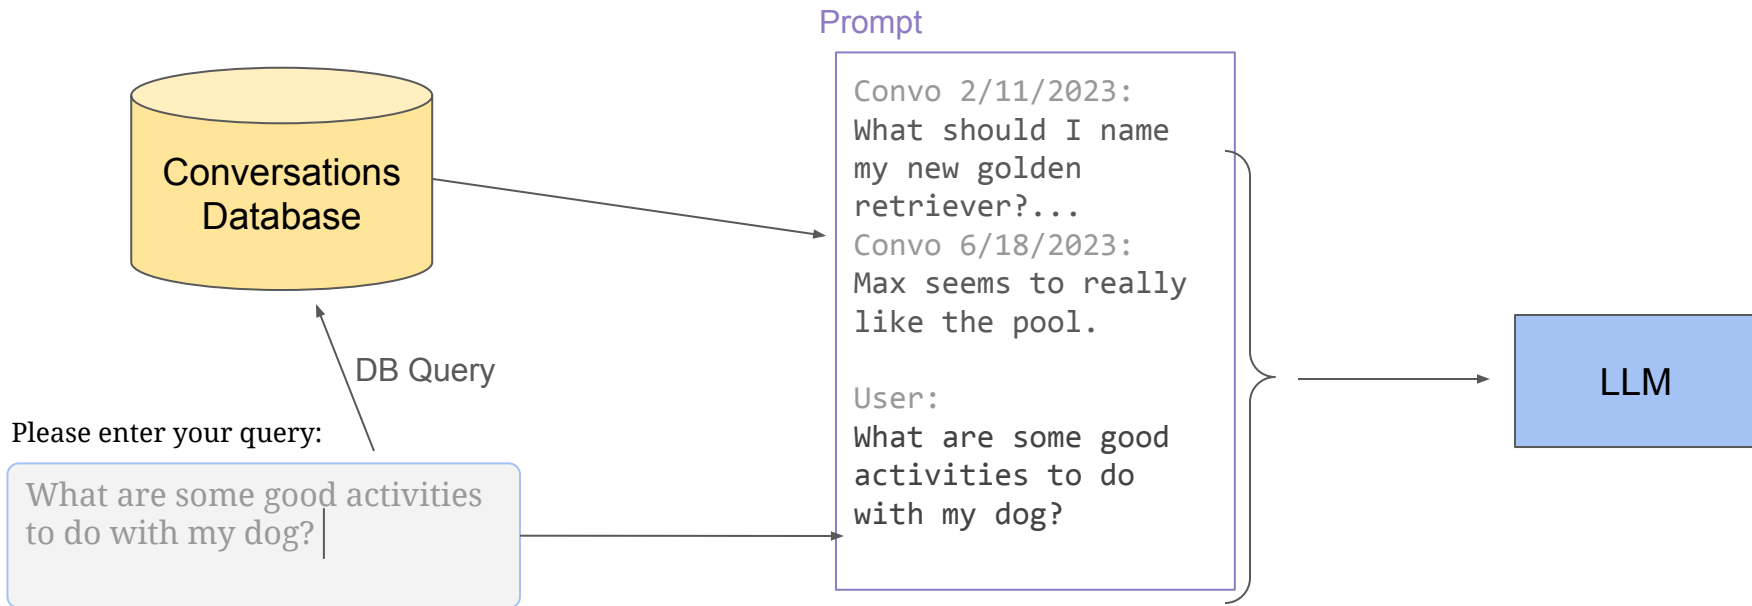
Although we cannot include all news and events (or all information from a target domain) in our prompt, we can include pieces of information that are relevant.

Combined with a software middleman (similar to shadow prompting), we can pull critical info into our prompt before it is processed:



Retrieval Augmented Generation (RAG)

A really common use-case of RAG is pulling in relevant information from past user conversations. This is how ChatGPT “knows about you”.



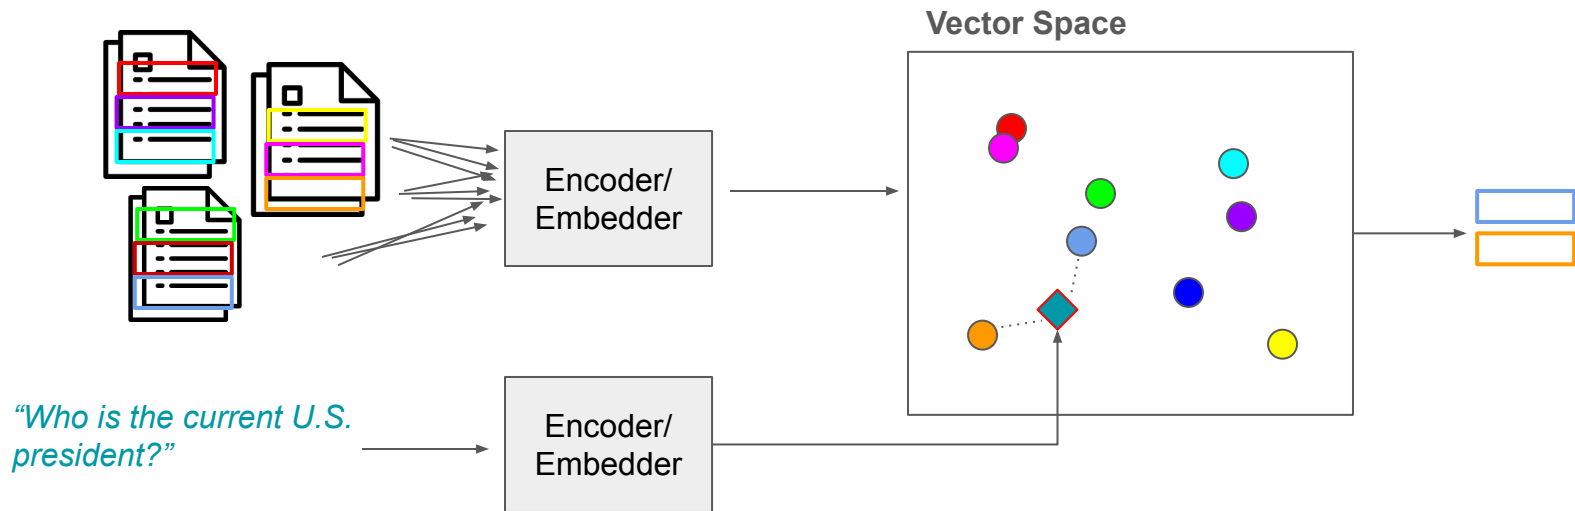
RAG

There are 4 basic steps:

- 1) Capture the user's query before sending to the LLM
- 2) Look up the most relevant pieces of information in a database
- 3) Construct a prompt that includes both the relevant information and the user's query
- 4) Feed all of this into the LLM

RAG Database Query

Typically the database for RAG is a vector store. Chunks of source documents are encoded as vectors in a database. The user's query is also encoded, and nearest-neighbors is used to find the N closest pieces of information.



Simple RAG Algorithm

Organize a database of documents or pieces of documents.

Using an embedding model (not a word embedding), embed all samples in the database.

When the user queries the system, embed the user query with the same embedding model.

Document

Embedding



0.1, 0.45, -0.55, 1.23...



-9.7, 0.15, -2.45, 1.54...



-3.7, 0.35, 2.45, -0.53...



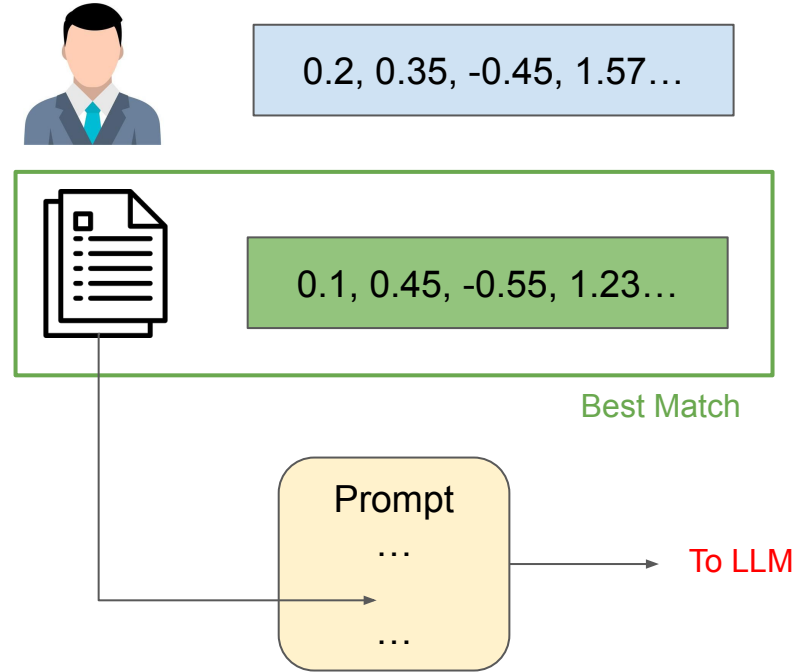
0.2, 0.35, -0.45, 1.57...

Simple RAG Algorithm (Cont.)

Select the documents with vector representations most similar to the user query.

- Small set of documents: use something like `torch.nn.functional.cosine_similarity`
- Many documents: tree-based data structure for fast lookup.

Copy the relevant text into a prompt, and proceed with LLM query.



LLM Inference Services + Tools

Inference Services

Large LLMs are often provided as an external service rather than run locally, due to size and cost.

In the preceding examples, invocations of an LLM would probably be API requests sent over the web, with the reply of the LLM coming back as a response.

Major LLM companies will provide APIs that you can hit from your software.



OpenAI



Claude

Gemini

Two Python Examples

```
from openai import OpenAI

# Initialize the OpenAI client.
# The API key is typically loaded from an environment variable (OPENAI_API_KEY)
client = OpenAI()

# Define the messages for the chat completion.
# The 'system' role provides instructions to the model.
# The 'user' role provides the user's input.
messages = [
    {"role": "system", "content": "You are a helpful assistant."},
    {"role": "user", "content": "Tell me a fun fact about space."},
]

try:
    # Create a chat completion request.
    # Specify the model to use (e.g., "gpt-3.5-turbo", "gpt-4").
    # Pass the defined messages.
    completion = client.chat.completions.create(
        model="gpt-3.5-turbo",
        messages=messages,
    )

    # Extract and print the model's response.
    # The response is typically found in the first choice of the completion.
    print(completion.choices[0].message.content)

except Exception as e:
    print(f"An error occurred: {e}")
```

```
import anthropic

# Initialize the Claude client with your API key
# It's recommended to set your API key as an environment variable (ANTHROPIC_API_KEY)
# For demonstration, you could also pass it directly: client = anthropic.Anthropic(api_key=...)
client = anthropic.Anthropic()

# Send a basic message to Claude
message = client.messages.create(
    model="claude-3-5-sonnet-20240620", # Specify the desired Claude model
    max_tokens=1024, # Set the maximum number of tokens
    messages=[
        {"role": "user", "content": "What is the capital of France?"}
    ],
)

# Print the response content
print(message.content)
```

Two Python Examples

```
from openai import OpenAI
```

```
# Initialize the OpenAI client.
```

```
# The API key is typically loaded from an environment variable (OPENAI_API_KEY)
```

```
client = OpenAI()
```

```
# Define the messages for the chat completion.
```

```
# The 'system' role provides instructions to the model.
```

```
# The 'user' role provides the user's input.
```

```
messages = [
    {"role": "system", "content": "You are a helpful assistant."},
    {"role": "user", "content": "Tell me a fun fact about space."},
]
```

```
try:
```

```
# Create a chat completion request.
```

```
# Specify the model to use (e.g., "gpt-3.5-turbo", "gpt-4").
```

```
# Pass the defined messages.
```

```
completion = client.chat.completions.create(
    model="gpt-3.5-turbo",
    messages=messages,
)
```

```
# Extract and print the model's response.
```

```
# The response is typically found in the first choice of the completion.
```

```
print(completion.choices[0].message.content)
```

```
except Exception as e:
```

```
    print(f"An error occurred: {e}")
```

```
import anthropic
```

```
# Initialize the Claude client with your API key
```

```
# It's recommended to set your API key as an environment variable (ANTHROPIC_API_KEY)
```

```
# For demonstration, you could also pass it directly: client = anthropic.Anthropic(api_key=API_KEY)
```

```
client = anthropic.Anthropic()
```

```
# Send a basic message to Claude
```

```
message = client.messages.create(
    model="claude-3-5-sonnet-20240620", # Specify the desired Claude model
    max_tokens=1024, # Set the maximum number of tokens to generate
```

```
    messages=[
        {"role": "user", "content": "What is the capital of France?"}
    ],
)
```

```
# Print the response content
```

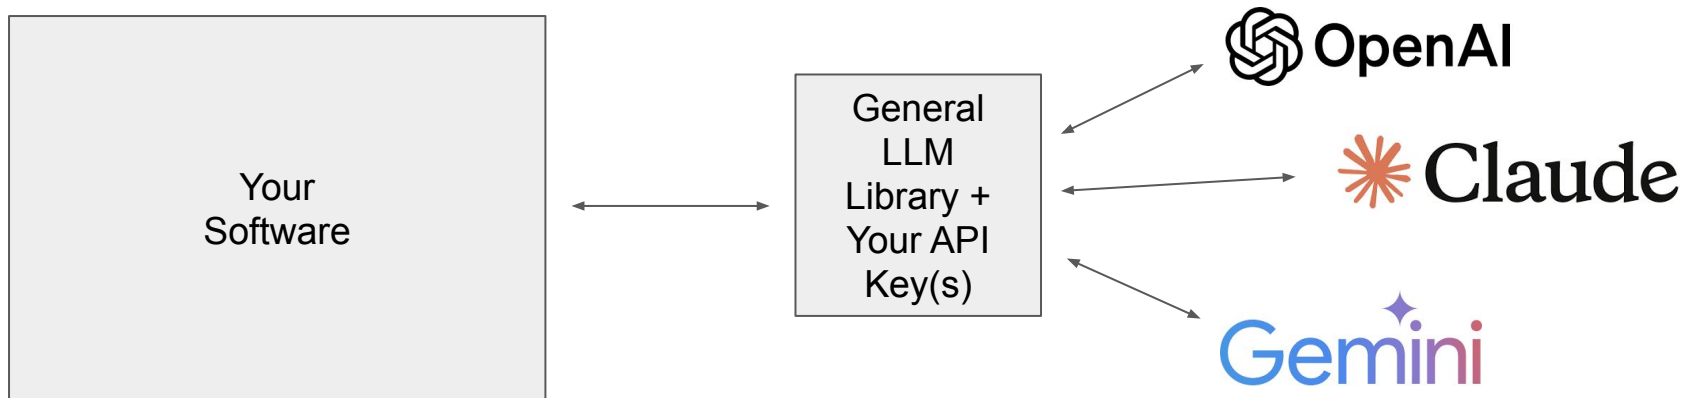
```
print(message.content)
```

Note the similar setup between these two providers. There has been a slow convergence towards standard API formats.

Libraries

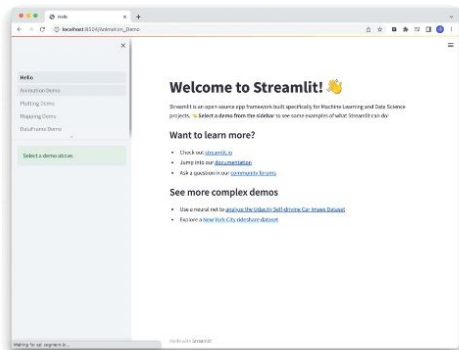
Additionally, many libraries exist which help handle the management of prompts, routing, etc and often abstract away any vendor-specific API formats. Popular options include **LangChain**, **DSPy**, and **Pydantic**.

Using these libraries, you typically supply your preferred LLM and the API key associated with that LLM, and the library backend will handle the rest.

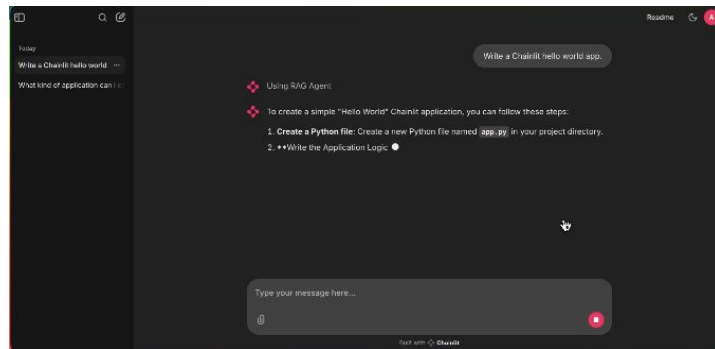


Browser UIs

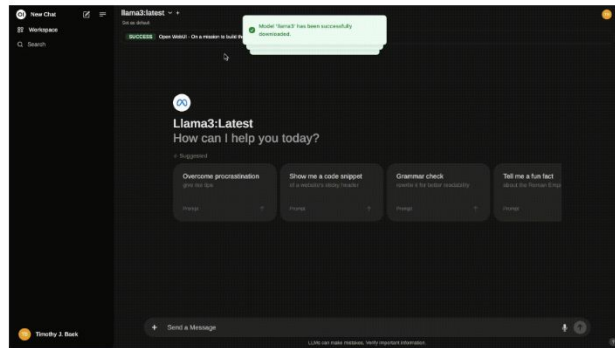
Alongside these libraries for communicating with LLMs, tools have also appeared to make nice UIs that go on top of them. These usually run in the browser and connect to python on the backend.



Streamlit



Chainlit



Open WebUI

Note / Opinion

General libraries like LangChain may become less useful as API formats converge.

LLM-agnostic libraries can also be incompatible with provider-specific offerings. Although, these will probably converge too.

Pydantic seems to be used by some LLM providers as the default system for packaging up messages, so I recommend using that.

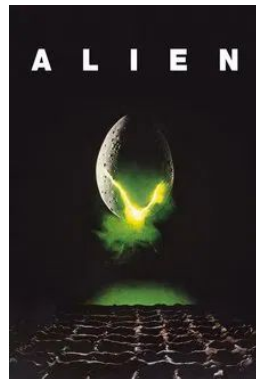
Hugging Face

For our assignment this week we will use Hugging Face for inference, because it comes with some free credits (a staggering 10 cents of credit per month).

We have already used Hugging Face here and there in our assignments. It has become the go-to for model and data hosting (which we will look at in a few slides), and lately is also pushing inference services.



Hugging Face is named after the hugging face emoji and uses it as their logo. It is not (like I thought) named after the creepy thing from alien.



Hugging Face Inference

To use HF Inference, you will need to an account (<https://huggingface.co/join>)

Just like other LLM providers, you will need to use a special key or token when calling their API in python. Instructions for this are in the appendix of this lecture.

```
import os
from openai import OpenAI # pip install openai

client = OpenAI(
    base_url="https://router.huggingface.co/v1",
    api_key=os.environ["HF_TOKEN"],
)

completion = client.chat.completions.create(
    model="Qwen/Qwen3-Next-80B-A3B-Instruct:novita",
    messages=[
        {
            "role": "user",
            "content": "What is the capital of France?"
        }
    ],
)

print(completion.choices[0].message)
```

Hugging Face Inference

To use HF Inference, you will need to an account (<https://huggingface.co/join>)

Just like other LLM providers, you will need to use a special key or token when calling their API in python. Instructions for this are in the appendix of this lecture.

Look familiar?

HF uses the OpenAI API to handle things, so again the format is the same.

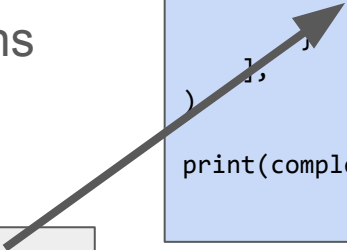
Note that we are not actually communicating with OpenAI's servers here. We are just using their python classes to package and transmit messages.

```
import os
from openai import OpenAI # pip install openai

client = OpenAI(
    base_url="https://router.huggingface.co/v1",
    api_key=os.environ["HF_TOKEN"],
)


completion = client.chat.completions.create(
    model="Qwen/Qwen3-Next-80B-A3B-Instruct:novita",
    messages=[
        {
            "role": "user",
            "content": "What is the capital of France?"
        }
    ],
)

print(completion.choices[0].message)
```



Inference Details

This is a very helpful link: <https://huggingface.co/inference/models>

 **Hugging Face**

[Models](#) [Datasets](#) [Spaces](#) [Community](#) [Docs](#) [Enterprise](#) [Pricing](#) [Learn more](#)

Inference Providers · Metrics for top trending models

Model	Provider	Status	Input \$/1M	Output \$/1M	Context	Latency (s)	Throughput (t/s)	Tools	Structured
Qwen/Qwen3-Next-80B-A3B-Instruct	novita	live	0.15	1.5	65536	0.74	154	Yes	No
Qwen/Qwen3-Next-80B-A3B-Instruct	together	live	0.15	1.5	262144	0.81	137	Yes	Yes
Qwen/Qwen3-Next-80B-A3B-Instruct	hyperbolic	live	-	-	262144	0.58	149	Yes	No
Qwen/Qwen3-Next-80B-A3B-Thinking	novita	live	0.15	1.5	65536	0.77	124	Yes	No
Qwen/Qwen3-Next-80B-A3B-Thinking	together	live	0.15	1.5	262144	0.76	187	Yes	Yes
Qwen/Qwen3-Next-80B-A3B-Thinking	hyperbolic	live	-	-	262144	0.56	139	No	No
openai/gpt-oss-20b	nebius	live	0.05	0.2	131072	0.31	185	Yes	Yes
openai/gpt-oss-20b	novita	live	0.05	0.2	131072	0.39	239	No	No
openai/gpt-oss-20b	together	live	0.05	0.2	131072	0.44	174	No	Yes
openai/gpt-oss-20b	fireworks-ai	live	0.05	0.2	131072	0.37	251	Yes	No
openai/gpt-oss-20b	groq	live	0.1	0.5	131072	0.14	986	Yes	No
openai/gpt-oss-20b	nscale	live	0.05	0.2	131072	0.41	122	Yes	Yes
openai/gpt-oss-20b	hyperbolic	live	0.1	0.1	131072	0.29	152	No	No
openai/gpt-oss-120b	nebius	live	0.15	0.6	131072	0.28	149	Yes	Yes
openai/gpt-oss-120b	cerebras	live	0.25	0.69	-	0.29	1074	Yes	No
openai/gpt-oss-120b	novita	live	0.1	0.5	131072	0.47	190	Yes	Yes

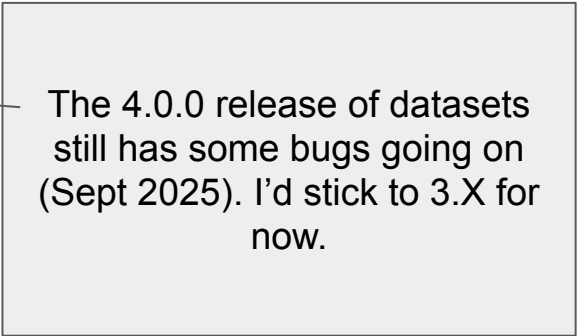
More Hugging Face

Main HF Libraries

HF has lots of offerings that are not tied to an account. Primarily, HF serves as a repository for open source models and datasets.

For this week's assignment we will be using:

- datasets (pip install datasets==3.1.0)
- transformers (pip install transformers)



The 4.0.0 release of datasets still has some bugs going on (Sept 2025). I'd stick to 3.X for now.

HF Datasets

Browsing

Downloading / using

Format of the dataset (train/test, columns, etc)

Process

Streaming

<https://huggingface.co/datasets?modality=modality:text&sort=trending>

Datasets

The screenshot shows the Hugging Face Datasets page. The top navigation bar includes the Hugging Face logo, a search bar, and links to Models, Datasets (highlighted with a red box), Spaces, Community, Docs, Enterprise, Pricing, and a user profile icon. The left sidebar contains filters for Modalities (3D, Audio, Document, Geospatial, Image, Tabular, Text, Time-series, Video), Size (rows) (a slider from <1K to >1T), and Format (json, csv, parquet, imagefolder, soundfolder, webdataset, text, arrow). The main content area displays a grid of datasets, each with a thumbnail icon, name, and metadata (Viewer, Updated, Size, Downloads, Likes). The datasets are sorted by trending.

Hugging Face Search models, datasets, users...

Models **Datasets** Spaces Community Docs Enterprise Pricing

Datasets 335,584 Filter by name Full-text search Sort: Trending

Dataset Name	Updated	Size	Downloads	Likes
HuggingFaceFW/finetpds	Updated 13 days ago	475M	74.6k	546
fka/awesome-chatgpt-prompts	Updated Jan 5	203	46.7k	9.09k
InternRobotics/OmniWorld	Updated about 6 hours ago	2.45M	9.93k	49
LucasFang/FLUX-Reason-6M	Updated 10 days ago	5.89M	34.5k	72
HuggingFaceM4/FineVision	Updated 17 days ago	24.2M	246k	336
Josephgflowers/Finance-Instruct-500k	Updated Mar 1	518k	3.13k	192
SamsungResearch/TRUEBench	Updated 6 days ago	19	385	18
JDhruv14/Bhagavad-Gita_Dataset	Updated 8 days ago	701	3.49k	50
lmms-lab/LLaVA-OneVision-1.5-Instruct-Data	Updated 3 days ago	12.9M	14.6k	14
InternRobotics/MotionMillion	Updated 3 days ago	1.25M	811	23
InternRobotics/InternData-M1	Updated 3 days ago	687k	1.56k	16
lmms-lab/LLaVA-One-Vision-1.5-Mid-Training-85M	Updated about 5 hours ago	18.1M	9.2k	12
jupyter-agent/jupyter-agent-dataset	Updated 11 days ago	95.8k	5.72k	140
SpatialVID/SpatialVID-HQ	Updated 5 days ago	365k	12.1k	22
m-a-p/DeepWriting-20K	Updated 13 days ago	35.8k	788	18
cais/mmlu	Updated Mar 8, 2024	231k	357k	544

Datasets

A huggingface dataset is a collection of samples (rows) where each sample has entries in various columns.

Datasets can be further organized into subsets, and are often split into a train and test division.

Datasets: allenai / c4 like 468 Follow Ai2 4.07k Dataset card

Subset (113)
en · ~365M rows (showing the first 2.57M)

Split (2)
train · ~365M rows (showing the first 2.21M)

Search this dataset

text	timestamp	url
string · lengths	string · lengths	string · lengths
<div><div></div><div>14184k</div></div>	<div><div></div><div>1919</div></div>	<div><div></div><div>142.53k</div></div>
Beginners BBQ Class Taking Place in Missoula! Do you want to get better at making delicious BBQ? Yo...	2019-04-25 12:57:54	https://klyq.com/beginners-bbq-class-in-missoula/
Discussion in 'Mac OS X Lion (10.7)' started by axboi87, Jan 20, 2012. I've got a 500gb internal...	2019-04-21 10:07:13	https://forums.macrumors.com/threads/larger-disk-to-smaller-disk.1311329/
Foil plaid lycra and spandex shortall with metallic slinky insets. Attached metallic elastic belt with...	2019-04-25 10:40:23	https://awishcometrue.com/Catalogs/C1/V1960-Find-A-Way
How many backlinks per day for new site? Discussion in 'Black Hat SEO' started by Omoplata, Dec 3,...	2019-04-21 12:46:19	https://www.blackhatworld.com/seo/how-backlinks-per-day-for-new-site.258615
The Denver Board of Education opened the 2017-18 school year with an update on projects that includ...	2019-04-20 14:33:21	http://bond.dpsk12.org/category/news/
BANGALORE CY JUNCTION SBC to GONDIA JUNCTION G train timings, routes, stops, and complete info. A...	2019-04-20 04:25:39	https://tatkalforsure.com/trains-betw-stations/bangalore-cy-junction-sbc-to
I thought I was going to finish the 3rd season of the Wire tonight. But there was a commentary on...	2019-04-18 14:16:05	https://karaokegal.livejournal.com/17
The rich get richer and the poor get poorer eh? Or is it the rich think different and play by a...	2019-04-23 00:39:43	http://www.iammeek.com/2018/06/the-ri-and-poor-get-poorer.html
Biomedics 1 Day Extra are daily replacement disposable contact lenses by CooperVision Hydron...	2019-04-26 09:38:13	https://www.webcontacts.com.au/Biomed-lenses/Biomedics-1-Day-Extra-90-pack

Datasets

A huggingface dataset is a collection of samples (rows) where each sample has entries in various columns.

Datasets can be further organized into subsets, and are often split into a train and test division.

This is the “english” subset, with 365M entries.

Datasets: allenai / c4 like 468 Follow Ai2 4.07k Dataset card

Subset (113)
en · ~365M rows (showing the first 2.57M)

Split (2)
train · ~365M rows (showing the first 2.21M)

Search this dataset

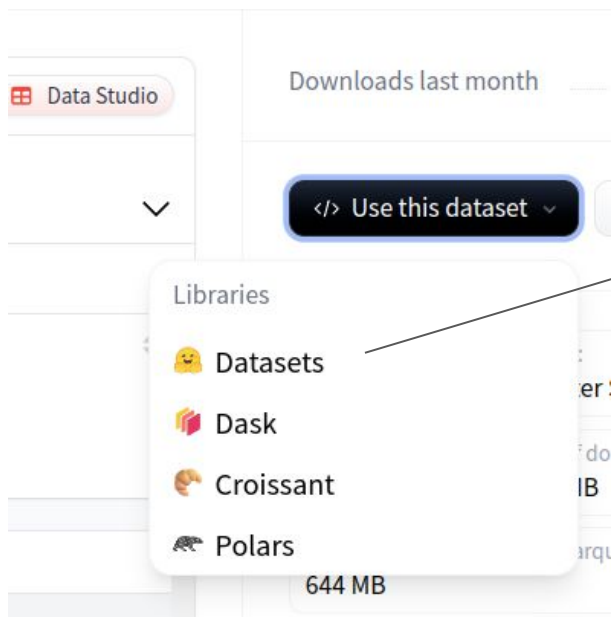
text string · lengths	timestamp string · lengths	url string · lengths
14 184k	19 19	14 2.53k
Beginners BBQ Class Taking Place in Missoula! Do you want to get better at making delicious BBQ? Yo...	2019-04-25 12:57:54	https://klyq.com/beginners-bbq-class-in-missoula/
Discussion in 'Mac OS X Lion (10.7)' started by axboi87, Jan 20, 2012. I've got a 500gb internal...	2019-04-21 10:07:13	https://forums.macrumors.com/threads/larger-disk-to-smaller-disk.1311329/
Foil plaid lycra and spandex shortall with metallic slinky insets. Attached metallic elastic belt with...	2019-04-25 10:40:23	https://awishcometrue.com/Catalogs/C1/V1960-Find-A-Way
How many backlinks per day for new site? Discussion in 'Black Hat SEO' started by Omoplata, Dec 3,...	2019-04-21 12:46:19	https://www.blackhatworld.com/seo/how-backlinks-per-day-for-new-site.258615
The Denver Board of Education opened the 2017-18 school year with an update on projects that includ...	2019-04-20 14:33:21	http://bond.dpsk12.org/category/news/
BANGALORE CY JUNCTION SBC to GONDIA JUNCTION G train timings, routes, stops, and complete info. A...	2019-04-20 04:25:39	https://tatkalforsure.com/trains-betw-stations/bangalore-cy-junction-sbc-to
I thought I was going to finish the 3rd season of the Wire tonight. But there was a commentary on...	2019-04-18 14:16:05	https://karaokegal.livejournal.com/17
The rich get richer and the poor get poorer eh? Or is it the rich think different and play by a...	2019-04-23 00:39:43	http://www.iammeek.com/2018/06/the-ri-and-poor-get-poorer.html
Biomedics 1 Day Extra are daily replacement disposable contact lenses by CooperVision Hydron...	2019-04-26 09:38:13	https://www.webcontacts.com.au/Biomed-lenses/Biomedics-1-Day-Extra-90-pack

This is the “train” split.

Three columns

Using a dataset in Python

Smaller datasets often have an example of loading with the datasets library:



```
How to use from the Datasets library
```

```
from datasets import load_dataset

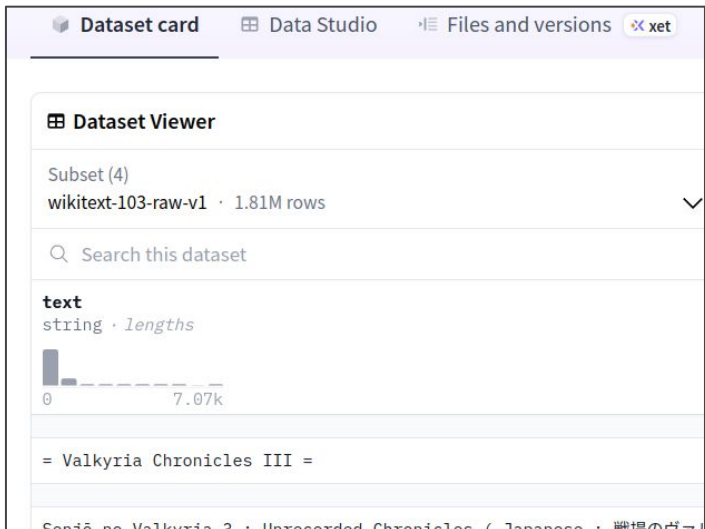
ds = load_dataset("Salesforce/wikitext", "wikitext-103-raw-v1")
```

Copy

This example is from
<https://huggingface.co/datasets/Salesforce/wikitext>

Using a dataset in python

Pasting the code into a script, we can examine the dataset. Note that it will need to download the first time you run.



```
from datasets import load_dataset

# loads the whole dataset
ds = load_dataset("Salesforce/wikitext",
                  "wikitext-2-raw-v1")

# look at the available columns
print(ds.column_names)
# prints: {'test': ['text'], 'train': ['text'],
#         'validation': ['text']}
# we can see there are three splits, each with only one
# column called "text".

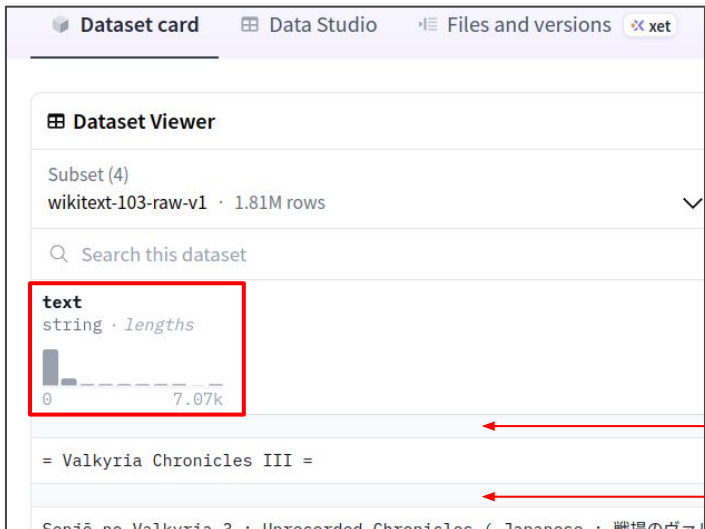
ds = ds["train"]
print("Samples:", len(ds))
# prints: Samples: 36718

# print out the first few entries
print(ds[0])
print(ds[1])
print(ds[2])
# prints:
# {'text': ''}
# {'text': ' = Valkyria Chronicles III = \n'}
# {'text': ''}

# The first and third rows are empty.
```

Using a dataset in python

Pasting the code into a script, we can examine the dataset. Note that it will need to download the first time you run.



```
from datasets import load_dataset

# loads the whole dataset
ds = load_dataset("Salesforce/wikitext",
                  "wikitext-2-raw-v1")

# look at the available columns
print(ds.column_names)
# prints: {'test': ['text'], 'train': ['text'],
#         'validation': ['text']}
# we can see there are three splits, each with only one
# column called "text".

ds = ds["train"]
print("Samples:", len(ds))
# prints: Samples: 36718

# print out the first few entries
print(ds[0])
print(ds[1])
print(ds[2])
# prints:
# {'text': ''}
# {'text': ' = Valkyria Chronicles III = \n'}
# {'text': ''}

# The first and third rows are empty.
```

This is a peculiarity of the dataset. We can confirm by looking on the HF page.

Processing the Dataset

There are all sorts of things you can do to refine or modify a dataset once it is loaded. See: <https://huggingface.co/docs/datasets/en/process>

For example, some really common operations might be things like:

- Remove columns that you do not need
- Shuffle the dataset order
- For testing purposes, only use the first N samples
- etc


What about that first dataset?

The initial example, c4, is huge (<https://huggingface.co/datasets/allenai/c4>). For this reason, there is no “Use this Dataset” button. However, there are options to download it directly.

If we scroll down the page, there are some examples anyway. What would happen if we used them?

```
from datasets import load_dataset

# English only
en = load_dataset("allenai/c4", "en")
```



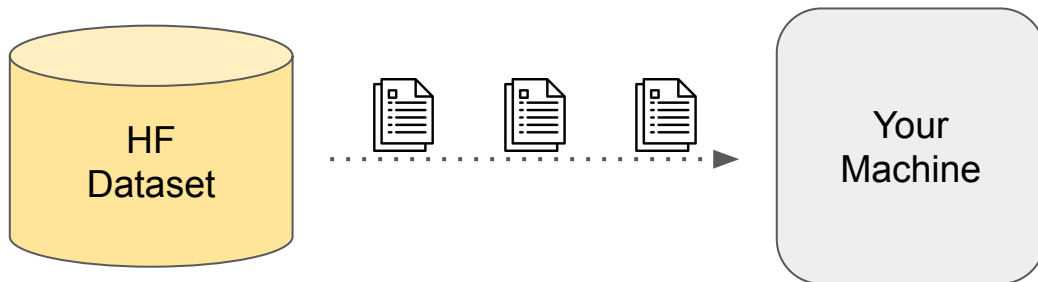
This command will download maybe 100GB to your computer. The entire dataset (all subsets) is 13GB.

Streaming

Hugging Face has an awesome alternative called streaming datasets which makes this much easier:

```
en = load_dataset("allenai/c4", "en", streaming=True)
```

This will create an **IterableDataset**, which is a dataset of undefined length that can be iterated. When your python script needs the next item from the dataset, it is downloaded on the fly (and deleted after use). This allows you to use giant datasets without ever storing them locally.



Streaming

If we set a dataset to stream, we can then iterate through it and samples are pulled as needed.

Note that this may start up a little slower than you expect, since the loader will pull down data in batches as well as supporting files.

```
from datasets import load_dataset

ds = load_dataset("allenai/c4", "en",
                  split="train", streaming=True)

# with an explicit iterator
it = iter(ds)
for i in range(10):
    sample = next(it)
    print(sample)

# with an implicit iterator
count = 0
for sample in ds:
    print(sample)
    count += 1
    if count==10:
        break
```

Streaming + Training?

Streaming can be used in a training loop, but it is not recommended.

There are a variety of practical issues with streaming when trying to train on a large amount of data:

- Download speed can easily be a bottleneck
- Need consistent internet connection
- Cannot shuffle data except for in a rolling pool
 - We can maintain a local buffer of N samples and sample randomly from this, but it is a coarse approximation to true shuffling.
- **IterableDataset** has no `len()`, which is needed to set up learning rate schedules.

HF Transformers

Hugging Face is also a repository for models. Most popular open-source models are available. Small-ish models can be downloaded and used locally via the **transformers** library (pip install transformers).

Note that this also includes relevant tokenizers. To process text, you need to load a model and its corresponding tokenizer.

Fortunately, this is very easy!

AutoModels and AutoTokenizers

HF Models and Tokenizers can be loaded by name. Similar to datasets, you will see a “Use this Model” button. For example, the GPT2 model (<https://huggingface.co/openai-community/gpt2>):

like 2.95k Follow OpenAI community 1.58k

PyTorch TensorFlow JAX LiteRT Rust ONNX Safetensors English

next-generation-inference License: mit

community 121

Train Deploy Use this model

Downloads last month
12,339,920

Safetensors

Model size 137M params Tensor type F32

Files info

Inference Providers NEW

```
# Load model directly
from transformers import AutoTokenizer, AutoModelForCausalLM

tokenizer = AutoTokenizer.from_pretrained("openai-community/gpt2")
model = AutoModelForCausalLM.from_pretrained("openai-community/gpt2")
```

AutoTokenizers

Calling `AutoTokenizer()` will return an object that can convert strings into text and vice versa.

You can call the object or use `encode()` to `tokenizer.encode` text, and `tokenizer.decode()` to convert back.

```
from transformers import AutoTokenizer,
AutoModelForCausalLM

tokenizer =
AutoTokenizer.from_pretrained("openai-community/gpt2")

text = "This is some example text."

tokens = tokenizer(text, return_tensors='pt')

# important
tokens = tokens["input_ids"][0]
print(tokens)

text_again = tokenizer.decode(tokens)
print(text_again)

# prints:
# tensor([1212,  318,  617, 1672, 2420,   13])
# This is some example text.
```

AutoTokenizers

Notes:

- We can pass “return_tensors” to select pytorch as the data scheme used
- The output of encoding is a dictionary with keys “input_ids” and “attention_mask”, which are used for the model forward pass.
- The values are given a batch dimension, so we also need to select the first “item in the batch”, hence the [0].

```
from transformers import AutoTokenizer,
AutoModelForCausalLM

tokenizer =
AutoTokenizer.from_pretrained("openai-community/gpt2")

text = "This is some example text."

tokens = tokenizer(text, return_tensors='pt')

# important
tokens = tokens["input_ids"][0]
print(tokens)

text_again = tokenizer.decode(tokens)
print(text_again)

# prints:
# tensor([1212,  318,  617, 1672, 2420,   13])
# This is some example text.
```

AutoModels

We can similarly create a model by name to run locally. Only do this for small models.

The reason the tokenizer output is an unwieldy dictionary is so that it can be directly passed into the model as arbitrary keyword arguments.

```
from transformers import AutoTokenizer,
AutoModelForCausalLM

tokens = tokenizer(text, return_tensors='pt')
model =
AutoModelForCausalLM.from_pretrained("openai-community/gpt
2")
output = model(**tokens)

print(output["logits"].shape)
# prints (1, 6, 50257)
# this is the raw output distribution at each position
```

Note that **output** is also a dictionary, here with fields “logits” and “past_key_values”. You should recognize these as the inputs to a sampler!

Review Assignment

Appendices

Access Token

The screenshot shows the Hugging Face user interface. At the top, the Hugging Face logo and a search bar are visible. The navigation bar includes links for Models, Datasets, Spaces, Community, and a user profile icon. The user profile icon is circled in red, and a red arrow points from it to a dropdown menu. The dropdown menu contains options for Profile, Notifications, New Model, New Dataset, New Space, New Collection, Create organization, Usage Quota, Private Storage, Zero GPU, Get Hugging Face PRO, Settings, Access Tokens, Billing, Changelog, and Sign Out. A red arrow points from the 'Access Tokens' option in the dropdown menu to the main content area. The main content area shows the user's profile information, including the name 'Ted' and the username 'estaleyjhu'. The 'Access Tokens' section is highlighted, and it states 'You have no Access Token'.

Hugging Face Search models, datasets, users...

Models Datasets Spaces Community Profile **estaleyjhu**

Ted
estaleyjhu

Switch

Profile

Account

Authentication

Organizations

Billing

Access Tokens

SSH and GPG Keys

Inference Providers

Webhooks

Papers

Notifications

Jobs **NEW**

Local Apps and Hardware

Gated Repositories

Access Tokens

User Access Tokens

Access tokens authenticate your identity to the Hugging Face Hub and allow applications to perform actions on your behalf. **Do not share your Access Tokens with anyone**; we regularly check for leaked Access Tokens and remove them immediately.

You have no Access Token

Access Token

Access Tokens

User Access Tokens



+ Create new token

Access tokens authenticate your identity to the Hugging Face Hub and allow applications to perform actions based on token permissions. ⓘ **Do not share your Access Tokens with anyone**; we regularly check for leaked Access Tokens and remove them immediately.

You have no Access Token

Access Token

Fine-grained

Token type

Fine-grained

Read

Write

! This cannot be changed after token creation.

Token name

token1

Name it something

User permissions (estaleyjhu)

Repositories

- ☐ Read access to contents of all repos under your personal namespace
- ☐ Read access to contents of all public gated repos you can access
- ☐ Write access to contents/settings of all repos under your personal namespace

Inference

Inference

- ☒ Make calls to Inference Providers
- ☐ Make calls to your Inference Endpoints ⓘ
- ☐ Manage your Inference Endpoints ⓘ

Access Token

Override any user-level or org-level permissions set below for the specified repositories. The token

Save your Access Token

×

Save your token value somewhere safe. **You will not be able to see it again after you close this modal.** If you lose it, you'll have to create a new one.

hf_

Copy

Name	Permissions
token1	FINEGRAINED

Done

Example Screenshot

```
import os
from openai import OpenAI

client = OpenAI(
    base_url="https://router.huggingface.co/v1",
    api_key=HF_TOKEN, # paste your key in here or load from a variable
)

completion = client.chat.completions.create(
    model="openai/gpt-oss-20b:hyperbolic",
    messages=[
        {
            "role": "user",
            "content": "List the first 10 digits of Pi in reverse order."
        }
    ],
)

# >>>>> To just get the reply:
print(completion.choices[0].message.content)
```

Note: This costs about \$0.00005, which is 0.05% of your free quota on HF. As long as you do not ask for a novel or stick this inside some sort of loop, you should be fine.